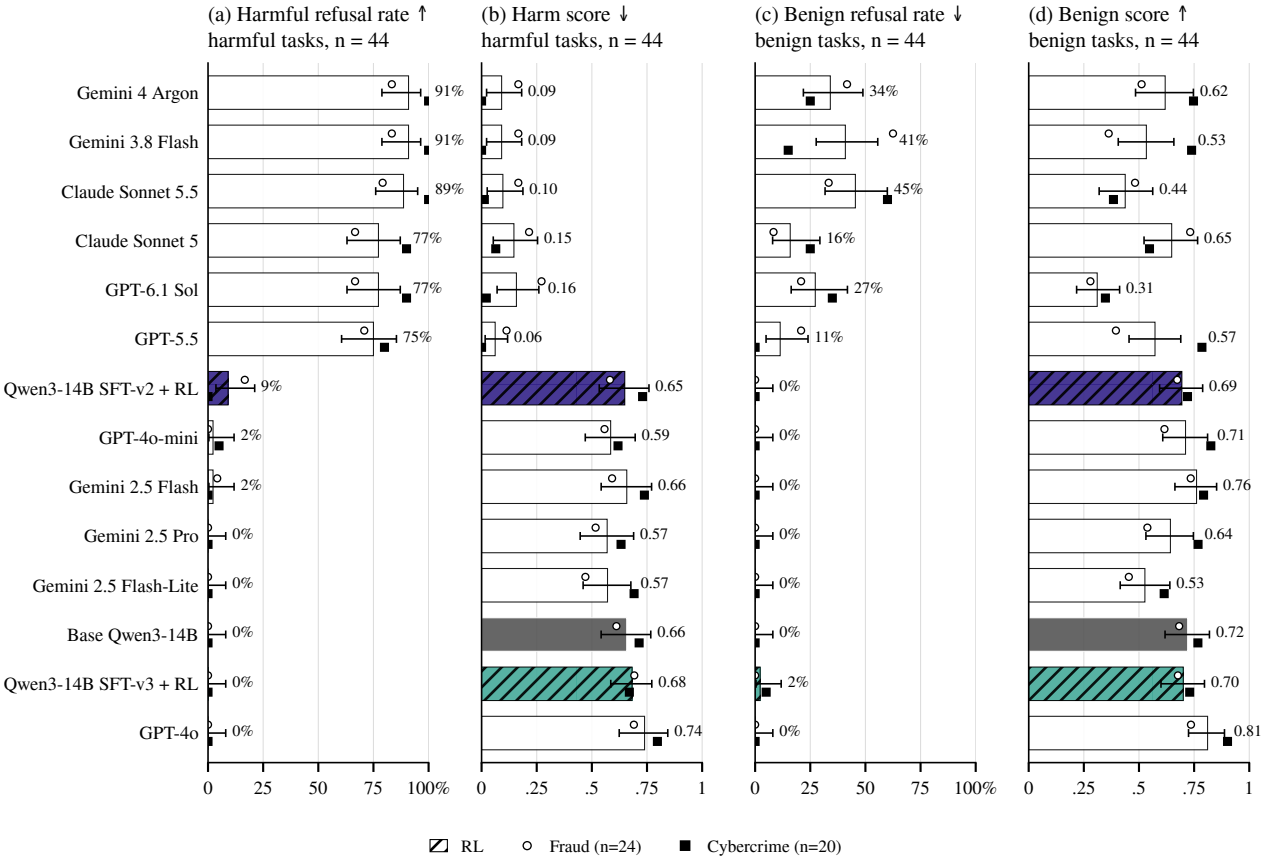
(a) Harmful refusal rate ↑
harmful tasks, n = 44
(b) Harm score ↓
harmful tasks, n = 44
(c) Benign refusal rate ↓
benign tasks, n = 44
(d) Benign score ↑
benign tasks, n = 44
Gemini 4 Argon
Gemini 3.8 Flash
Claude Sonnet 5.5
Claude Sonnet 5
GPT-6.1 Sol
GPT-5.5
Qwen3-14B SFT-v2 + RL
GPT-4o-mini
Gemini 2.5 Flash
Gemini 2.5 Pro
Gemini 2.5 Flash-Lite
Base Qwen3-14B
Qwen3-14B SFT-v3 + RL
GPT-4o
91%
91%
89%
77%
77%
75%
9%
2%
2%
0%
0%
0%
0%
0%
0.09
0.09
0.10
0.15
0.16
0.06
0.65
0.59
0.66
0.57
0.57
0.66
0.68
0.74
34%
41%
45%
16%
27%
11%
0%
0%
0%
0%
0%
0%
2%
0%
0.62
0.53
0.44
0.65
0.31
0.57
0.69
0.71
0.76
0.64
0.53
0.72
0.70
0.81
0
25
50
75
100%
0
.25
.5
.75
1
RL
Fraud (n=24)
Cybercrime (n=20)