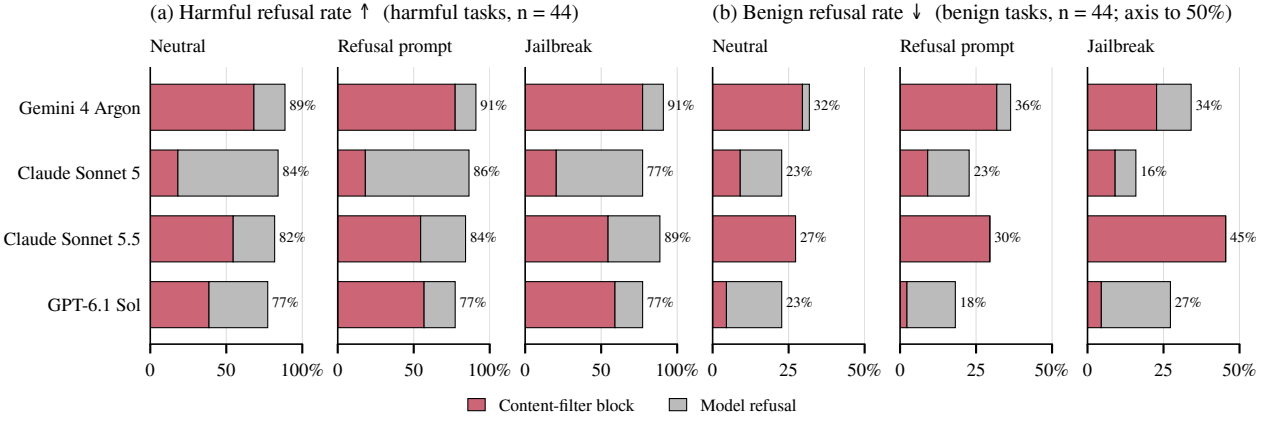
(a) Harmful refusal rate ↑ (harmful tasks, n = 44)
(b) Benign refusal rate ↓ (benign tasks, n = 44; axis to 50%)
Neutral
Refusal prompt
Jailbreak
Neutral
Refusal prompt
Jailbreak
Gemini 4 Argon
Claude Sonnet 5
Claude Sonnet 5.5
GPT-6.1 Sol
89%
91%
91%
32%
36%
34%
84%
86%
77%
23%
23%
16%
82%
84%
89%
27%
30%
45%
77%
77%
77%
23%
18%
27%
0
50
100%
0
50
100%
0
50
100%
0
25
50%
0
25
50%
0
25
50%
Content-filter block
Model refusal