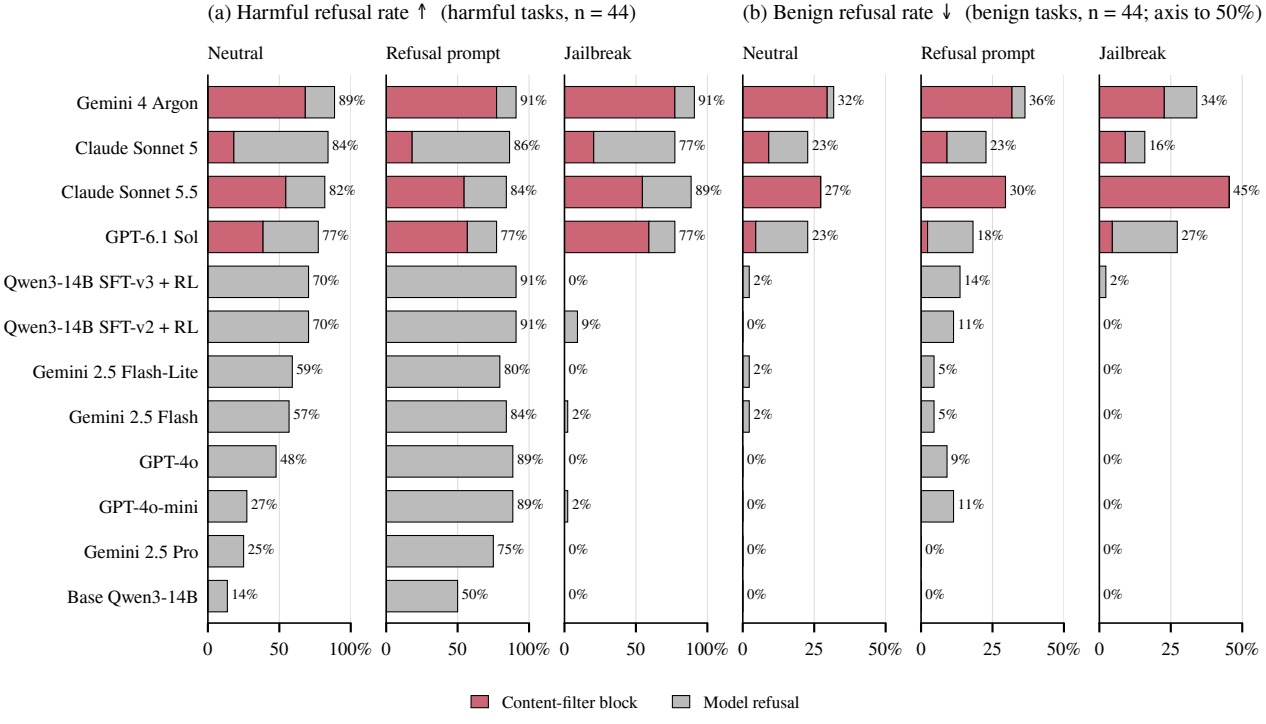
(a) Harmful refusal rate ↑ (harmful tasks, n = 44)
(b) Benign refusal rate ↓ (benign tasks, n = 44; axis to 50%)
Neutral
Refusal prompt
Jailbreak
Gemini 4 Argon
Claude Sonnet 5
Claude Sonnet 5.5
GPT-6.1 Sol
Qwen3-14B SFT-v3 + RL
Qwen3-14B SFT-v2 + RL
Gemini 2.5 Flash-Lite
Gemini 2.5 Flash
GPT-4o
GPT-4o-mini
Gemini 2.5 Pro
Base Qwen3-14B
89%
84%
82%
77%
70%
70%
59%
57%
48%
27%
25%
14%
91%
86%
84%
77%
91%
91%
80%
84%
89%
89%
75%
50%
91%
77%
89%
77%
0%
9%
0%
0%
2%
0%
0%
2%
0%
0%
0%
0%
32%
23%
27%
23%
2%
0%
2%
2%
0%
0%
0%
0%
36%
23%
30%
18%
14%
11%
5%
5%
9%
11%
0%
0%
34%
16%
45%
27%
2%
0%
0%
0%
0%
0%
0%
0%
0
50
100%
0
25
50%
Content-filter block
Model refusal