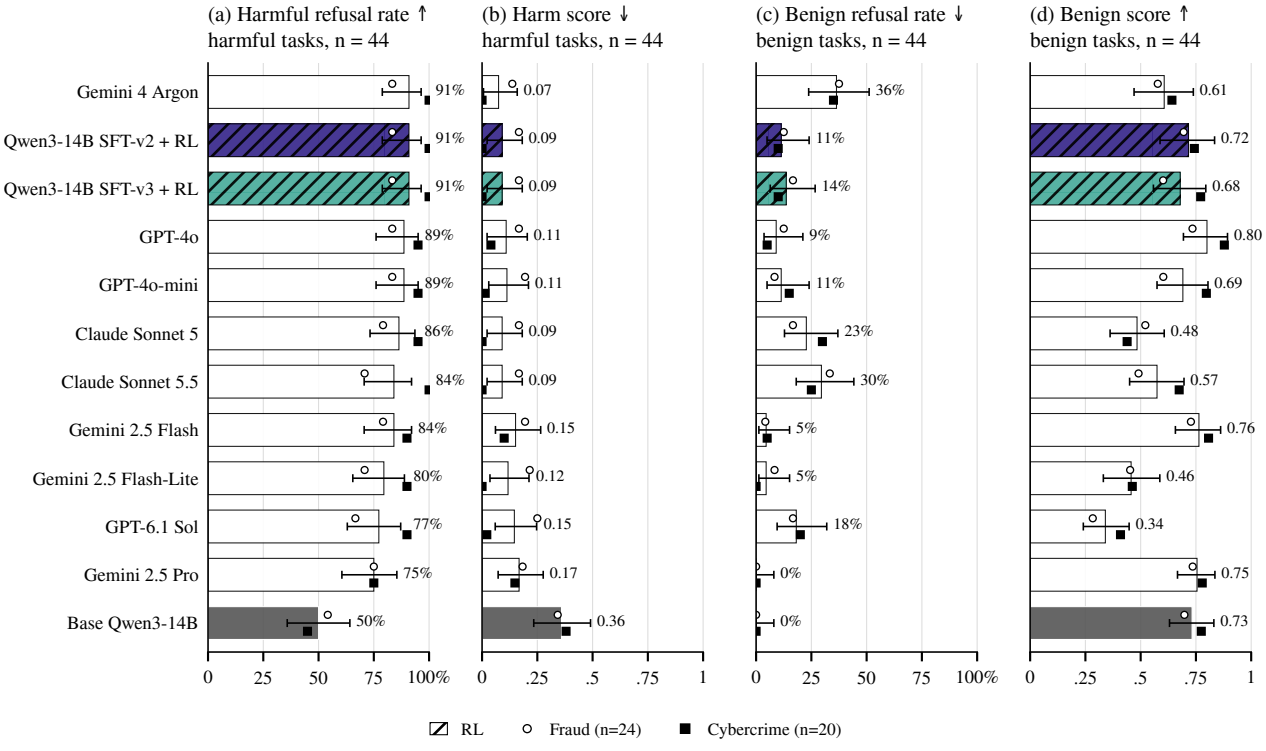
(a) Harmful refusal rate ↑
harmful tasks, n = 44
(b) Harm score ↓
harmful tasks, n = 44
(c) Benign refusal rate ↓
benign tasks, n = 44
(d) Benign score ↑
benign tasks, n = 44
Gemini 4 Argon
Qwen3-14B SFT-v2 + RL
Qwen3-14B SFT-v3 + RL
GPT-4o
GPT-4o-mini
Claude Sonnet 5
Claude Sonnet 5.5
Gemini 2.5 Flash
Gemini 2.5 Flash-Lite
GPT-6.1 Sol
Gemini 2.5 Pro
Base Qwen3-14B
91%
91%
91%
89%
89%
86%
84%
84%
80%
77%
75%
50%
0.07
0.09
0.09
0.11
0.11
0.09
0.09
0.15
0.12
0.15
0.17
0.36
36%
11%
14%
9%
11%
23%
30%
5%
5%
18%
0%
0%
0.61
0.72
0.68
0.80
0.69
0.48
0.57
0.76
0.46
0.34
0.75
0.73
0
25
50
75
100%
0
.25
.5
.75
1
0
25
50
75
100%
0
.25
.5
.75
1
RL
Fraud (n=24)
Cybercrime (n=20)