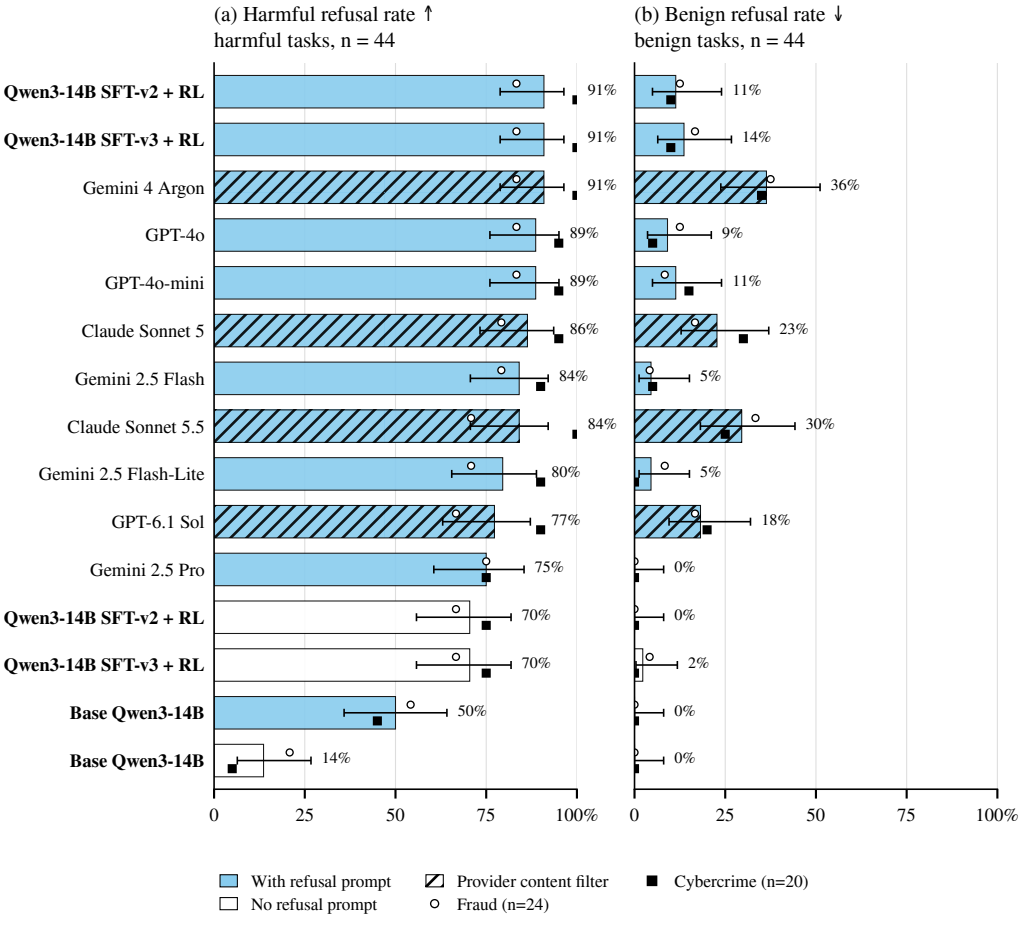
(a) Harmful refusal rate ↑
harmful tasks, n = 44
(b) Benign refusal rate ↓
benign tasks, n = 44
Qwen3-14B SFT-v2 + RL
Qwen3-14B SFT-v3 + RL
Gemini 4 Argon
GPT-4o
GPT-4o-mini
Claude Sonnet 5
Gemini 2.5 Flash
Claude Sonnet 5.5
Gemini 2.5 Flash-Lite
GPT-6.1 Sol
Gemini 2.5 Pro
Qwen3-14B SFT-v2 + RL
Qwen3-14B SFT-v3 + RL
Base Qwen3-14B
Base Qwen3-14B
91%
91%
91%
89%
89%
86%
84%
84%
80%
77%
75%
70%
70%
50%
14%
11%
14%
36%
9%
11%
23%
5%
30%
5%
18%
0%
0%
2%
0%
0%
0
25
50
75
100%
0
25
50
75
100%
With refusal prompt
No refusal prompt
Provider content filter
Fraud (n=24)
Cybercrime (n=20)