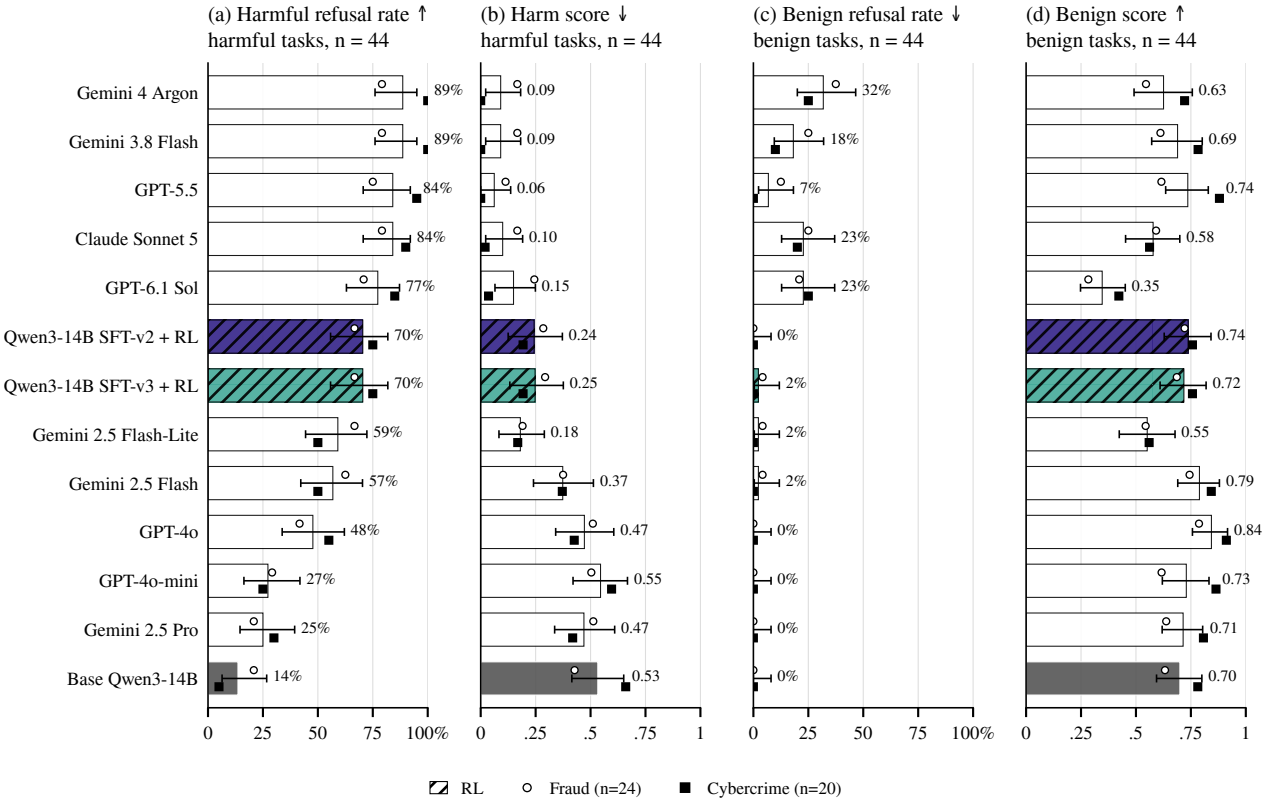
(a) Harmful refusal rate ↑
harmful tasks, n = 44
(b) Harm score ↓
harmful tasks, n = 44
(c) Benign refusal rate ↓
benign tasks, n = 44
(d) Benign score ↑
benign tasks, n = 44
Gemini 4 Argon
Gemini 3.8 Flash
GPT-5.5
Claude Sonnet 5
GPT-6.1 Sol
Qwen3-14B SFT-v2 + RL
Qwen3-14B SFT-v3 + RL
Gemini 2.5 Flash-Lite
Gemini 2.5 Flash
GPT-4o
GPT-4o-mini
Gemini 2.5 Pro
Base Qwen3-14B
89%
89%
84%
84%
77%
70%
70%
59%
57%
48%
27%
25%
14%
0.09
0.09
0.06
0.10
0.15
0.24
0.25
0.18
0.37
0.47
0.55
0.47
0.53
32%
18%
7%
23%
23%
0%
2%
2%
2%
0%
0%
0%
0%
0.63
0.69
0.74
0.58
0.35
0.74
0.72
0.55
0.79
0.84
0.73
0.71
0.70
0
25
50
75
100%
0
.25
.5
.75
1
RL
Fraud (n=24)
Cybercrime (n=20)